

Architectural Design & Engineering Report

Autonomous Drone Security Analyst Agent

Author: Om Bhavsar — Candidate for AI Engineer Role at FlytBase

Date: August 2026

Assignment: FlytBase AI Engineer Assignment — Multimodal AI, Agentic Workflows & Cross-Domain Indexing

Repository: Private GitHub Repository (Added `assignments@flytbase.com` as collaborator)

Demo Video Links: Attached / Provided via Google Drive Link

1. Executive Summary

Autonomous docked drones represent a paradigm shift for industrial property perimeter surveillance. However, continuous aerial surveillance generates unmanageable volumes of unstructured video data, overwhelming human operators and causing alert fatigue through frequent false alarms.

This report presents the architectural design, engineering trade-offs, and production implementation of the **Autonomous Drone Security Analyst Agent**, an end-to-end multimodal AI platform that:

- Synchronizes Telemetry & Video:** Ingests high-frequency telemetry (GPS, altitude, heading, battery) alongside video frames while enforcing strict idempotency and clock drift validation.
 - Dense Multimodal Understanding:** Leverages a Vision-Language Model (VLM) abstraction that extracts strongly-typed `StructuredObservation` schemas (`objects`, `activities`, `risk\_indicators`, and `confidence`) rather than unconstrained strings.
 - Durable Cross-Domain Indexing:** Replaces volatile in-memory storage with **SQLite** as the single source of truth and an **FTS5 virtual table** with BM25 ranking for sub-millisecond lexical search.
 - Stateful Cross-Frame Entity Tracking:** Correlates multi-frame entity fingerprints over time, profiling trajectories (e.g. tracking a Blue Ford F150 arriving at 12:00 at the Garage and departing at 14:30 at the Main Gate).
 - Decoupled Risk & Alert Engine:** Mathematically separates visual detection confidence ($C \in [0, 1]$) from situational threat severity ($R \in [0, 1]$) to minimize false-positive guard dispatches.
 - Dual-Agent Conversational Analyst:** Implements both a high-performance fast-path agent and a **LangChain enterprise toolkit** with `@tool` integrations for conversational forensics and automated shift summarization.
 - Interactive SOC Dashboard & OpenCV Engine:** Features a dark-mode **NiceGUI** operations center and **OpenCV** visual frame renderer with real-time HUD overlays.
-

2. Problem Definition & Assumptions

A. The Operational Challenge

Industrial facilities (logistics hubs, energy sub-stations, manufacturing yards) face two core surveillance dilemmas:

- **High Ingestion Volume:** Continuous video feeds cannot be manually reviewed in real time.
- **Context Blindness:** Standard vision detectors trigger in isolation without historical awareness (e.g., flagging a delivery truck on every pass or alerting on daytime employees).


B. Core Architectural Assumptions

- Docked Drone Flight Cycle:** The drone operates on scheduled perimeter sweeps and automated responder missions, streaming structured telemetry (GPS, altitude, speed) and optical video frames.
- Edge-to-Cloud Decoupling:** Drone hardware performs initial stream capture and OpenCV HUD stamping; dense VLM extraction, stateful event tracking, and FTS5 indexing run on the surveillance station/server.
- Pluggable VLM Model Selection:** Production systems require both offline deterministic testing (Mock VLM) and live cloud multimodal inference (Gemini / GPT-4o-mini).
- Data Privacy:** Storing structured observations, bounding attributes, and telemetry in SQLite ensures audit compliance and low storage overhead without retaining massive petabytes of raw uncompressed video.

3. Feature Specification & Value Proposition

- **Value Proposition:** Maximizes industrial asset protection by converting raw aerial drone footage into a queryable semantic timeline, eliminating false-positive guard dispatches through context-aware alerting thresholds.
 - **Key Requirements:**
 - *Real-Time Synchronized Ingestion:* Telemetry-frame alignment, boundary validation, and clock drift detection.
 - *Structured Multimodal Logging:* Granular object classification, visual attribute parsing, and activity logging.
 - *Contextual Alerting:* Night loitering rules, multi-entry vehicle recurrence profiling, and restricted zone monitoring.
 - *Cross-Domain Hybrid Indexing:* Sub-millisecond text search across visual descriptions and metadata.
 - *Conversational Analyst Agent:* Tool-invoking natural language interface for security investigations and shift summaries.
-

4. System Architecture & Data Flow



5. Telemetry–Frame Synchronization & Idempotency

In real-world drone operations, network jitter and packet retransmissions can cause duplicated frames or clock drift.

- **Idempotency Gate:** Every incoming frame possesses a deterministic `frame\_id` (e.g. `FRM-20260818-0001`). Re-transmitted packets are caught by `FrameRepository` and skipped, preventing duplicate database entries.
- **Clock Drift & Boundary Checks:** `IngestionSynchronizer` verifies GPS coordinates within legal geographical boundaries ($-90^\circ \leq \text{Lat} \leq 90^\circ$, $-180^\circ \leq \text{Lon} \leq 180^\circ$), ensures non-negative altitude ($\geq 0\text{m}$), and warns if telemetry and frame capture timestamps drift by $>5.0\text{s}$.

6. VLM Selection Matrix: CLIP vs BLIP vs Generative VLMs

Choosing the optimal vision backbone requires balancing latency, operational cost, and semantic reasoning:

Architectural Conclusion:

We decoupled the VLM behind `BaseVLMEngine`. In production, a Generative VLM generates dense JSON observations, while the storage layer abstracts retrieval so that vector embeddings (CLIP) can be added seamlessly without breaking the pipeline.

7. Structured Observation Schema

Rather than relying on unstructured text descriptions, our VLM pipeline enforces a strongly-typed Pydantic schema:

```
{
  "description": "A blue Ford F150 pickup truck backing into the primary delivery garage space.",
  "objects": [
    {
      "object_type": "vehicle",
      "label": "Ford F150",
      "attributes": {"color": "blue", "type": "pickup truck"},
      "confidence": 0.96
    }
  ],
  "activities": [
    {
      "activity_type": "entering",
      "confidence": 0.93,
      "notes": "Maneuvering into loading bay at Garage Loading Bay"
    }
  ],
  "risk_indicators": ["vehicle_loading_access"],
  "vlm_confidence": 0.94
}
```

8. Cross-Frame Context & Stateful Event Tracking

Single-frame analysis cannot differentiate between a routine authorized entry and an unauthorized reconnaissance sweep.

The `EventTracker` builds stateful entity profiles:

- **Fingerprinting:** Normalized composite keys:

$\text{Key}_{\text{vehicle}} = \text{vehicle}:\text{label}:\text{color}$

$\text{Key}_{\text{person}} = \text{person}:\text{label}:\text{clothing}$

- **State Aggregation:**

- `occurrence_count`: Tracked across continuous flights.
 - `locations`: Ordered list of visited zones (e.g. `Garage Loading Bay -> Main Gate`).
 - `frame_ids`: Complete forensic audit trail linking every historical frame.
 - `activities_observed`: Historical behaviors across the property.
-

9. Decoupled Risk & Alert Engine

Differentiating Detection Confidence from Threat Severity

A core requirement for production systems is distinguishing:

- **Detection Confidence** ($\$C \in [0, 1]\$$): How certain the VLM is that an object or activity exists in the frame.
- **Threat Risk Score** ($\$R \in [0, 1]\$$): How dangerous that event is to property security.

$$\text{Risk Score} = \text{Clamp}_{[0,1]}(\text{BaseThreat} + W_{\text{night}} + W_{\text{zone}} + W_{\text{loitering}} + W_{\text{recurrence}} + \sum W_{\text{indicators}})$$

Severity Scale:

- $\$0.00 \leq R < 0.30 \implies \text{LOW}\$$
- $\$0.30 \leq R < 0.60 \implies \text{MEDIUM}\$$
- $\$0.60 \leq R < 0.80 \implies \text{HIGH}\$$
- $\$0.80 \leq R \leq 1.00 \implies \text{CRITICAL}\$$

Implemented Rules:

1. **NightLoiteringRule**: Triggers when an individual is detected loitering/pacing between 22:00 and 05:00 near restricted gates $\implies \text{CRITICAL}\$$ alert ($\$R=0.85\$$).
2. **RepeatedVehicleRule**: Triggers when an identified vehicle reaches $\geq 2\$$ sightings across property zones $\implies \text{MEDIUM}\$$ alert with full trajectory history.

10. Cross-Domain Indexing: Why SQLite + FTS5?

Rather than introducing heavy standalone vector database containers for simple keyword lookups, **SQLite + FTS5** was chosen:

1. **Durable Single-Node Persistence**: ACID compliance guarantees zero data loss upon application restart.
2. **Native Full-Text Search (FTS5)**: Provides Porter stemming, token normalization, and BM25 ranking.
3. **Sub-Millisecond Search Latency**: In-process memory mapping yields queries in $< 2\text{ms}\$$.
4. **Hybrid Retrieval**: Allows seamless SQL composition combining FTS text matching with relational filters (`'location_tag = 'Main Gate' AND severity = 'HIGH'`).
5. **Clean Abstraction**: The `HybridRetriever` interface isolates search, allowing a vector backend to be introduced without modifying business logic.

11. Conversational Security Analyst & LangChain Integration

The system supports a dual-agent conversational architecture:

- **Native Fast-Path Agent** (`app/agent/assistant.py`): Sub-5ms deterministic tool execution for real-time operations.
- **LangChain Enterprise Toolkit** (`app/agent/langchain_agent.py`): Exposes domain repositories as standard LangChain `@tool` functions (`search_surveillance_frames`, `get_security_alerts`, `get_cross_frame_events`, `get_mission_shift_summary`).

12. Computer Vision & OpenCV (cv2) Pipeline

Implemented in `app/pipeline/cv_utils.py`:

- **Telemetry HUD Overlays**: Stamps real-time heads-up displays directly onto video frames (GPS, Altitude, Heading, Battery, Time).
 - **Bounding Box Annotation**: Draws color-coded visual bounding boxes around detected entities (Vehicles, Pedestrians, Wildlife) and threat level banners.
 - **Optical Frame Storage**: Saves rendered, annotated JPEG frames to `data/rendered_frames/` for visual forensic review.
-

13. Concrete Results & Verification

Sample Outputs:

1. **Logs:** Frame `FRM-20260818-1200` @ Garage Loading Bay: "A blue Ford F150 pickup truck backing into the primary delivery garage space."
 2. **Alert:** Frame `FRM-20260818-0001` @ Main Gate: "[CRITICAL] Security Alert: Adult Individual in dark hoodie detected loitering after-hours at Main Gate (Altitude: 15.0m)."
 3. **Indexed Query:** `POST /search {"query": "truck"}` returned 2 matching frames with BM25 rank score `0.6521`.
 4. **Conversational Agent:** `POST /chat {"prompt": "How many times was the blue Ford truck observed today?"}` \$implies\$ "The Ford F150 (blue) was observed 2 time(s) today. First seen at 12:00 at Garage Loading Bay, last seen at 14:30 at Main Gate."
-

14. What Could Be Done Better Without Time Constraints

1. **On-Edge VLM Quantization:** Deploy 4-bit quantized Moondream2 or SmolVLM directly on drone companion computers (e.g. NVIDIA Jetson Orin Nano).
 2. **Live RTSP / WebRTC Stream Ingestion:** Connect directly to drone video feeds using OpenCV GStreamer pipelines.
 3. **Dense Vector Embeddings:** Combine `sentence-transformers` (`all-MiniLM-L6-v2`) with FTS5 for full hybrid lexical-dense vector retrieval.
 4. **Geo-Fencing Polygon Mapping:** Enable dynamic polygonal perimeter drawing on the NiceGUI map for custom restricted zone definitions.
-

15. AI-Assisted Co-Working & Engineering Log

In compliance with FlytBase's evaluation criteria, the project was engineered co-working with leading AI development tools:

16. Submission & Setup Summary

- **GitHub Repository:** Private repository created with code in `drone\_security\_agent`.
- **Contributor:** `assignments@flytbase.com` added as collaborator.
- **Run Instructions:**

```
cd drone_security_agent
pip install -r requirements.txt
python -m pytest           # Runs 21 passing tests
python dashboard.py        # Starts NiceGUI SOC on http://localhost:8080
uvicorn app.main:app --reload # Starts REST API on http://localhost:8000/docs
```

17. Conclusion

The **Autonomous Drone Security Analyst Agent** provides an enterprise-ready multimodal prototype that addresses every FlytBase evaluation benchmark:

- **Correctness & Performance:** Dense structured VLM outputs with multi-factor risk alerting.
- **Reasoning & Scalability:** Cross-frame entity trajectory tracking and decoupled architecture.
- **Innovation & Indexing:** SQLite FTS5 hybrid search, NiceGUI SOC dashboard, OpenCV visual renderer, and LangChain toolkit.
- **Documentation & Code Quality:** 21 passing automated tests, interactive Jupyter sandbox, 3-5 min video demo script, and comprehensive architectural justification.